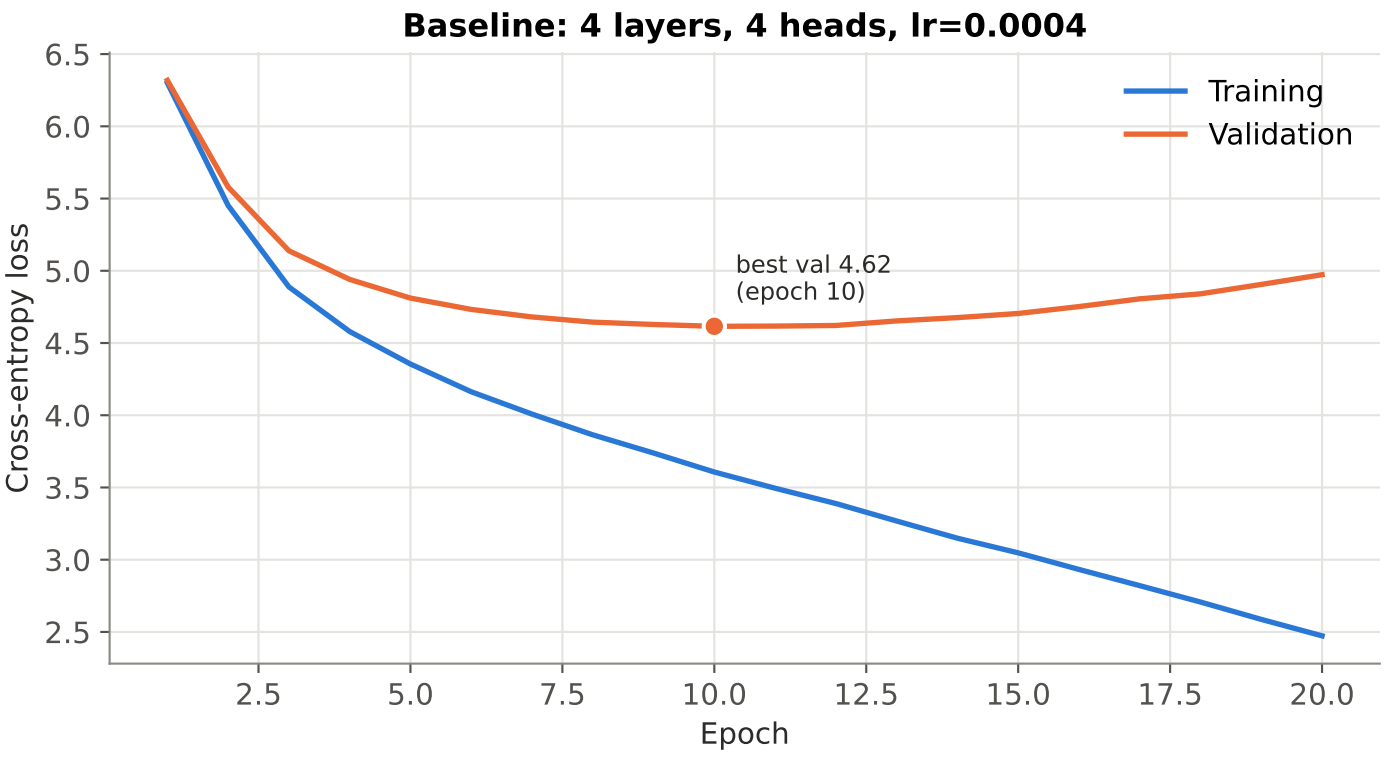
Baseline: 4 layers, 4 heads, lr=0.0004
Training
Validation
best val 4.62
(epoch 10)
Cross-entropy loss
Epoch
6.5
6.0
5.5
5.0
4.5
4.0
3.5
3.0
2.5
2.5
5.0
7.5
10.0
12.5
15.0
17.5
20.0