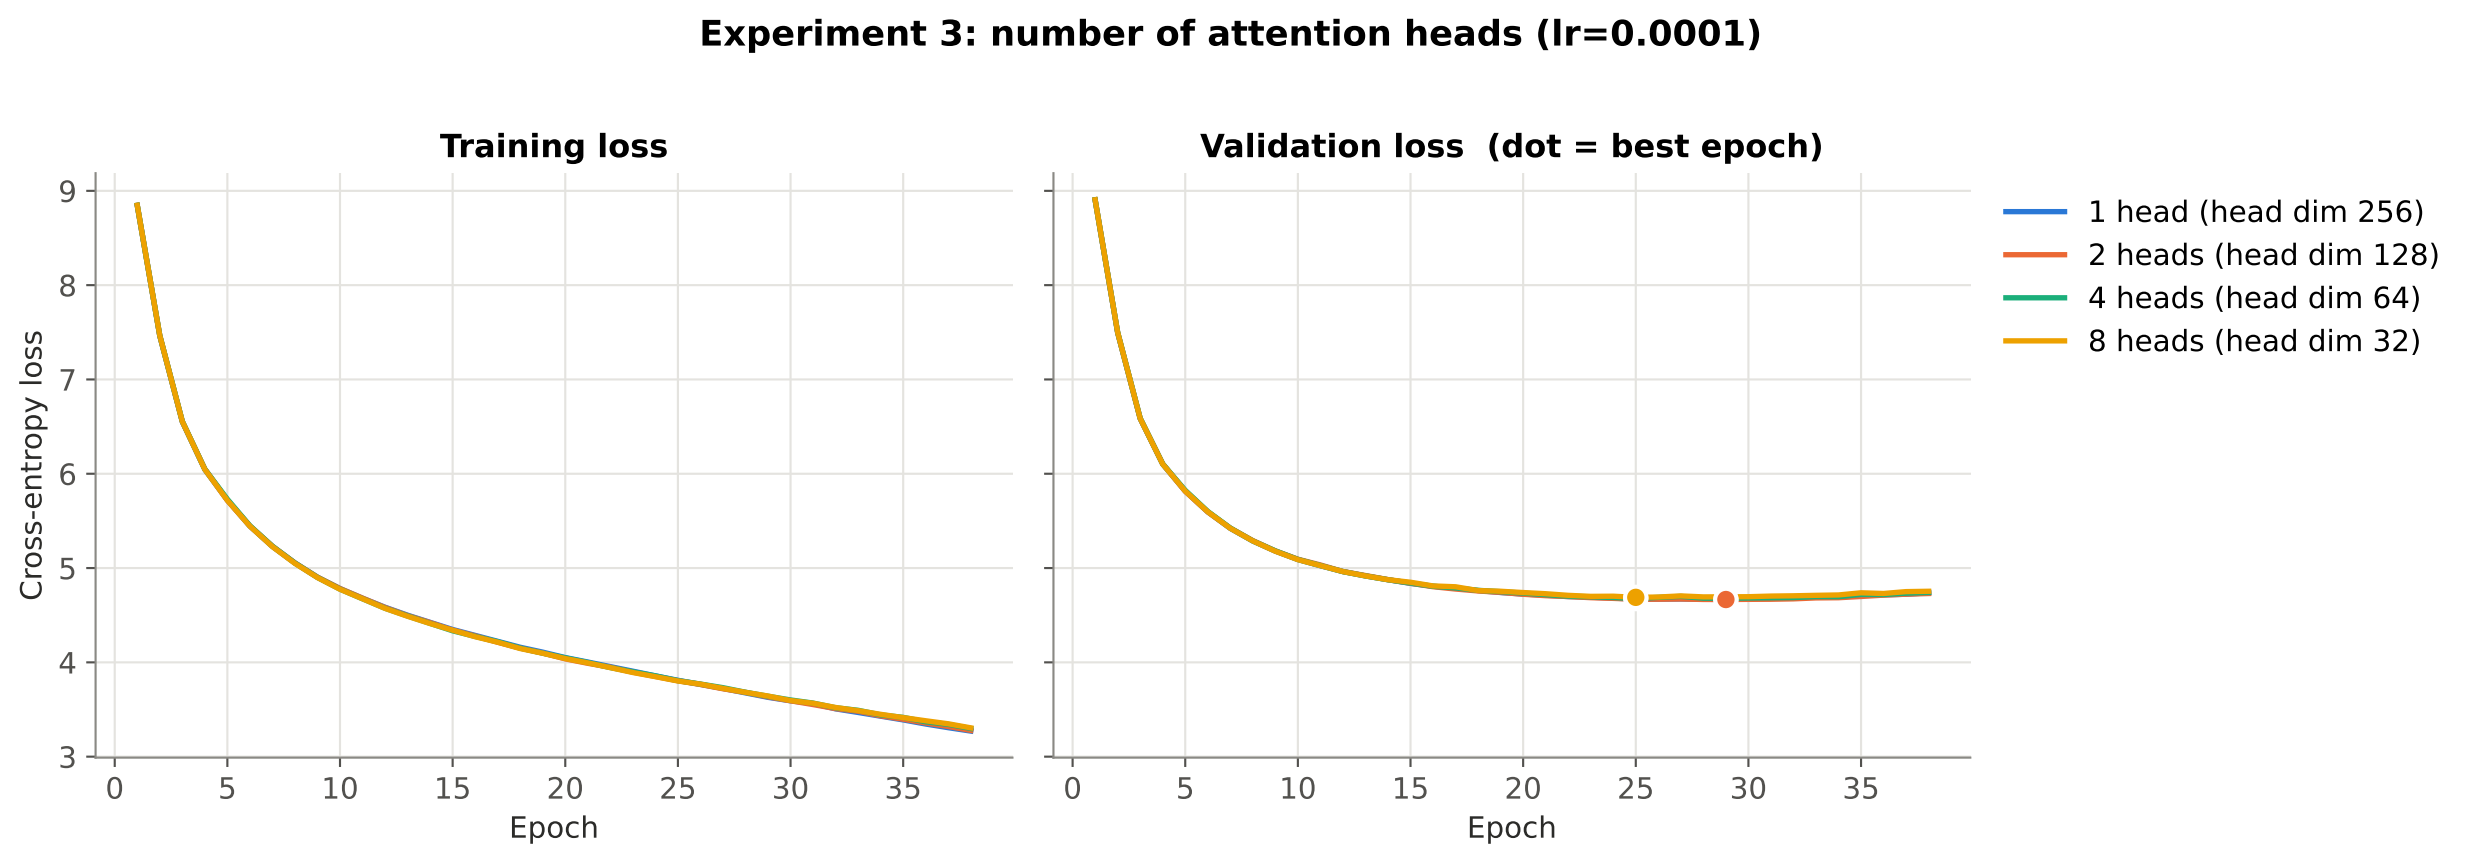
Experiment 3: number of attention heads (lr=0.0001)
Training loss
Validation loss (dot = best epoch)
Cross-entropy loss
Epoch
Epoch
1 head (head dim 256)
2 heads (head dim 128)
4 heads (head dim 64)
8 heads (head dim 32)