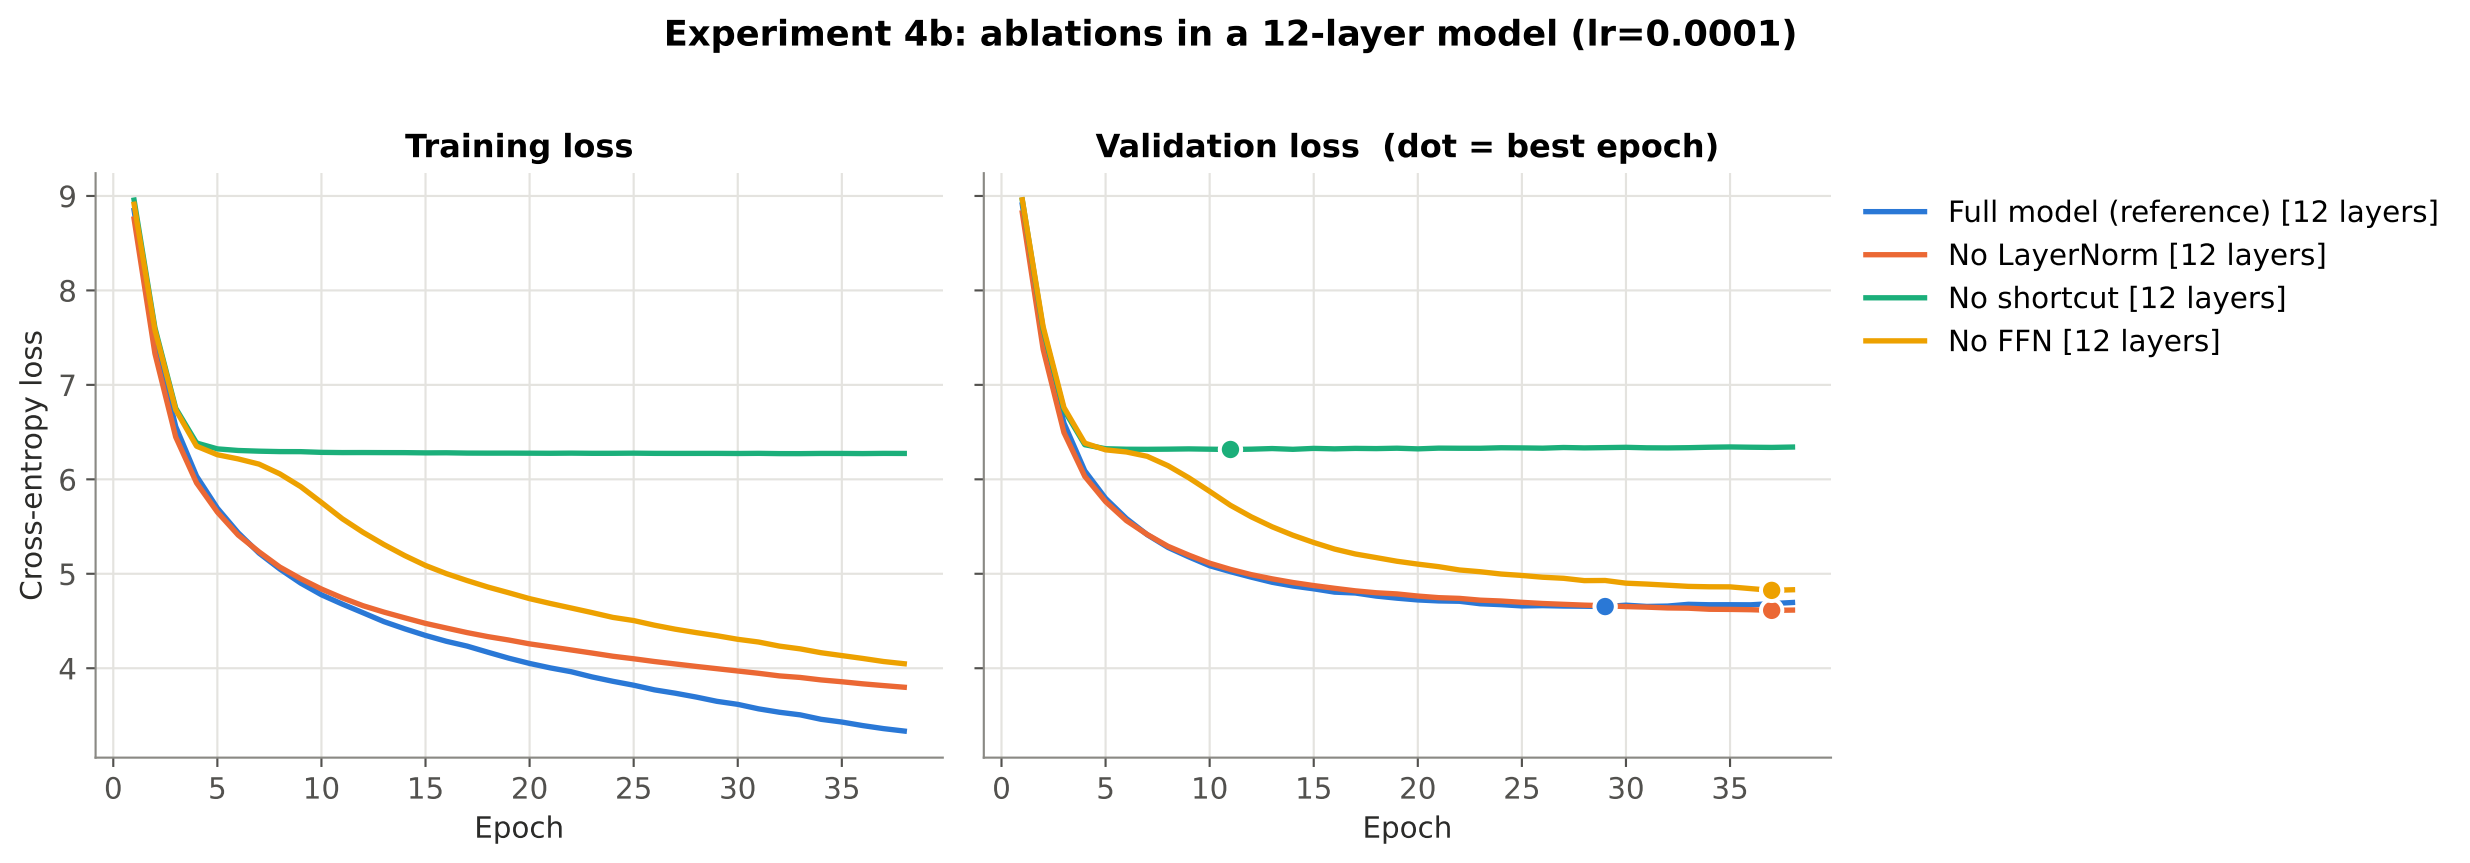
Experiment 4b: ablations in a 12-layer model (lr=0.0001)
Training loss
Validation loss (dot = best epoch)
Cross-entropy loss
Epoch
Epoch
Full model (reference) [12 layers]
No LayerNorm [12 layers]
No shortcut [12 layers]
No FFN [12 layers]