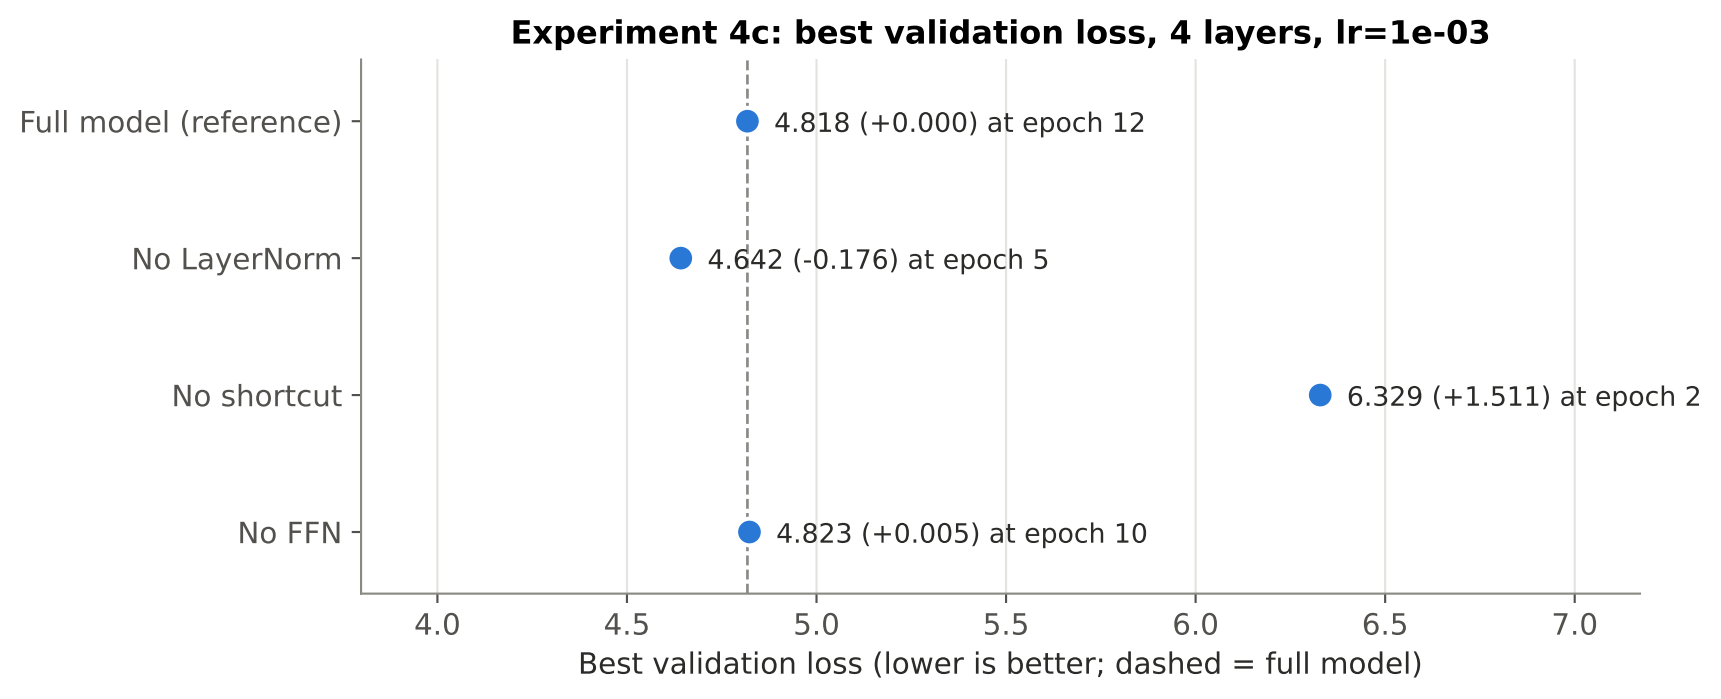
Experiment 4c: best validation loss, 4 layers, lr=1e-03
Full model (reference)
No LayerNorm
No shortcut
No FFN
4.818 (+0.000) at epoch 12
4.642 (-0.176) at epoch 5
6.329 (+1.511) at epoch 2
4.823 (+0.005) at epoch 10
4.0
4.5
5.0
5.5
6.0
6.5
7.0
Best validation loss (lower is better; dashed = full model)