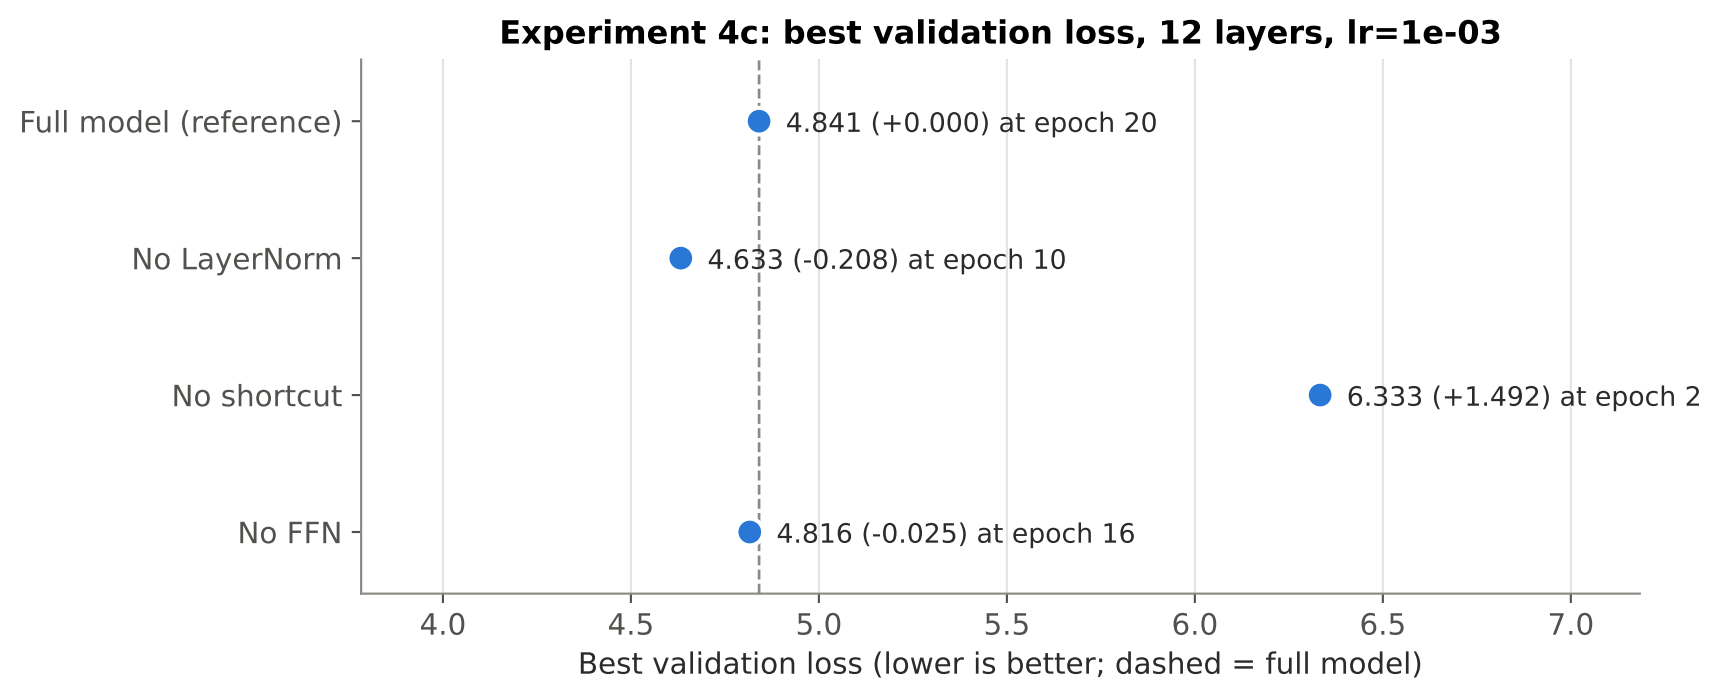
Experiment 4c: best validation loss, 12 layers, lr=1e-03
Full model (reference)
No LayerNorm
No shortcut
No FFN
4.841 (+0.000) at epoch 20
4.633 (-0.208) at epoch 10
6.333 (+1.492) at epoch 2
4.816 (-0.025) at epoch 16
4.0
4.5
5.0
5.5
6.0
6.5
7.0
Best validation loss (lower is better; dashed = full model)